



Using MySQL 5.6 Global Transaction IDs in Production

Stéphane Combaudon
August 27th, 2014

Agenda

- Introduction to GTIDs
- Daily DBA tasks
- Typical issues
- High availability solutions with GTIDs

Introduction to GTIDs

What is a GTID? (1)

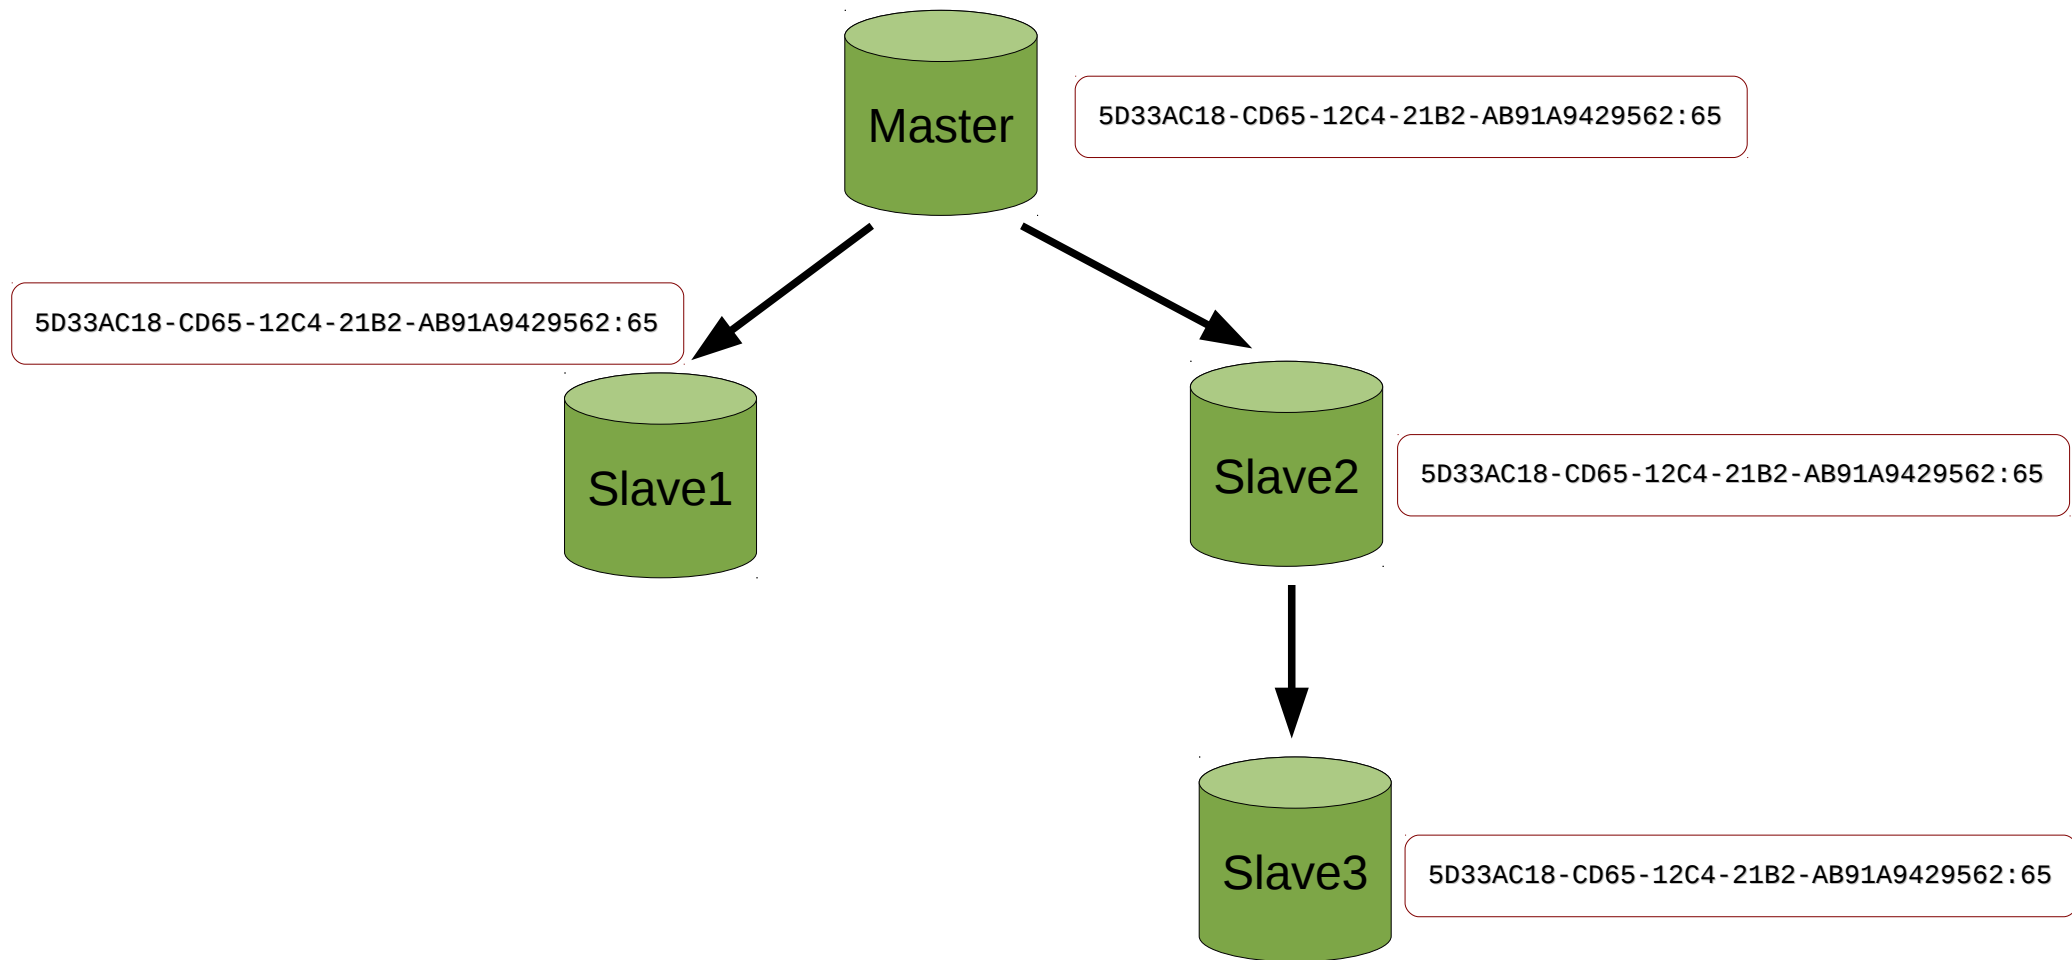
- Unique identifier of a transaction across all servers of a replication setup
- Available from MySQL 5.6
- Main goal: being able to change the replication topology easily
- MariaDB 10 has another implementation that is not compatible
 - Work is in progress to allow replication from MySQL 5.6 to MariaDB 10

What is a GTID? (2)

- 2 parts
 - `source_id:transaction_id`
 - `3E11FA47-71CA-11E1-9E33-C80AA9429562:1`
- Mapping to actual binlog file/position is kept in memory
- A sequence of GTIDs
 - `source_id:trx_start-trx_stop`
 - `3E11FA47-71CA-11E1-9E33-C80AA9429562:1-5`

Finding the position of an event

- Now that's easy, same for all servers!!



Transaction ordering

- Transaction counter is per instance, not global
- Say we have `xxx:101` and `xxx:102`, which came first?
 - `xxx:101`
- Say we have `xxx:101` and `yyy:102`, which came first?
 - We don't know

Enabling GTIDs

- Add following settings in my.cnf on all servers
 - `gtid_mode = ON`
 - `log_bin`
 - `log-slave-updates`
 - `enforce-gtid-consistency`
- Then restart all servers at the same time
 - Yes, it's mandatory :(
 - Will be improved in 5.7
 - Booking.com and Facebook have patches for online migration

Using GTID replication

- Once GTIDs are enabled on all servers, run
 - `CHANGE MASTER TO ..., MASTER_AUTO_POSITION = 1`
- `MASTER_LOG_POS` and `MASTER_LOG_FILE` are no longer needed

Replication protocol

- When slave connects to the master
 - Position-based replication
 - Master sends all transactions from the given offset
 - GTID-based replication
 - Sends the range of GTIDs it has executed
 - Master sends back all other transactions
 - Rule: a transaction with a given GTID can only execute once
- More on that new replication protocol later

Daily DBA tasks

Provisioning a slave

- mysqldump
 - `--master-data` now includes GTID information
 - Reload the dump and run `CHANGE MASTER TO ... MASTER_AUTO_POSITION=1`
- Percona XtraBackup
 - `xtrabackup_binlog_info` contains GTID information
 - After moving the backup, run `SET GLOBAL gtid_purged="XYZ"`
 - Then run `CHANGE MASTER TO ... MASTER_AUTO_POSITION=1`

Checking replication status (1)

- New columns for SHOW SLAVE STATUS

```
Retrieved_Gtid_Set: 41631daf-0295-11e4-9909-94dbc999324d:4-7  
Executed_Gtid_Set: 41631daf-0295-11e4-9909-94dbc999324d:1-7  
Auto_Position: 1
```

- Retrieved_Gtid_Set: GTIDs received by the slave, cleared after a server restart
- Executed_Gtid_Set: List of GTIDs executed. Here last executed transaction is 41631daf-0295-11e4-9909-94dbc999324d:7
- Auto_position: 1 if GTID-based replication is in use

Checking replication status (2)

- `gtid_executed`
 - Set of executed GTIDs, identical to `Executed_Gtid_Set`

```
Executed_Gtid_Set: 41631daf-0295-11e4-9909-94dbc999324d:1-7
Auto_Position: 1
1 row in set (0,00 sec)

sb2> show global variables like 'gtid_executed';
+-----+
| Variable_name | Value                                     |
+-----+
| gtid_executed | 41631daf-0295-11e4-9909-94dbc999324d:1-7 |
+-----+
```

- After several failovers, can be more complex

```
Retrieved_Gtid_Set: 4162896e-0295-11e4-9909-94dbc999324d:1-2
Executed_Gtid_Set: 4162896e-0295-11e4-9909-94dbc999324d:1-2,
41631daf-0295-11e4-9909-94dbc999324d:1-7,
4163bec4-0295-11e4-9909-94dbc999324d:1-2
Auto_Position: 1
```

Skipping transactions (1)

- `sql_skip_slave_counter = N` no longer works
 - Because of the new replication protocol, the transaction would automatically come back
 - It throws an error if you try to use it anyway
- But there's a solution!
 - Execute a fake trx with the GTID you want to skip
 - New replication protocol makes sure the real trx will not be executed

Skipping transactions (2)

- How to skip transaction `xxxx:nn`?
 - `STOP SLAVE;`
 - `SET gtid_next = 'xxxx:nn';`
 - `BEGIN;COMMIT;` # Fake transaction!
 - `SET gtid_next=automatic;`
 - `START SLAVE;`

Typical issues

Errant transactions

- What if you execute a trx locally on a slave?
 - It generates its own GTID
 - If the slave is promoted, trx is sent to all the servers
- That can bite on failover
 - Trx is not desired: sorry, now it is everywhere
 - Trx is no longer in the binlogs: sorry, this triggers a replication error

Detecting errant transactions

- Executed_Gtid_Set of any slave should always be a subset of Executed_Gtid_Set of master
 - GTID_SUBSET() can be used for this check
- If `SELECT GTID_SUBSET(slave_set,master_set)` returns 0, you have errant transactions
- Then use `GTID_SUBTRACT(master_set,slave_set)` to identify them

Fixing errant transactions

- Inject an empty transaction on all other servers of the topology
- If you have to run local transactions, use SET `sql_log_bin = 0`

Holes in the GTID sequence

- Holes are not allowed, but there are bugs
 - <http://bugs.mysql.com/bug.php?id=71575>
 - <http://bugs.mysql.com/bug.php?id=71376> (fixed in 5.6.18)
- That can lead to issues similar to those hit with errant transactions
- No tool is currently checking holes

I/O performance issues

- `log_bin + log_slave_updates` adds some I/O overhead
- Mapping between GTID and actual position is kept in memory
 - On initial connect, dump threads has to reverse scan the master's binlogs
 - This can be expensive if the slave is far behind

High Availability solutions with GTIDs

MySQL Utilities

- Set of Python scripts to ease administration of MySQL servers
- Free and open source, developed by Oracle
- <http://dev.mysql.com/doc/workbench/en/mysql-utilities.html>

Overview of mysqlfailover

- Health monitoring and automatic failover
 - Target topology: 1 master, N slaves
- A few MySQL settings are required
 - `--log-slave-updates, --enforce-gtid-consistency, gtid_mode = ON`
 - `--report-host, --report-port`
 - `--master-info-repository=TABLE`

Existing modes

- Elect
 - Chooses a candidate from a list. If none can be promoted, exits with an error
- Auto (default)
 - Same as elect, but if no candidate is suitable, any other slave can be promoted
- Fail
 - Perform health monitoring, exits with an error if the master fails

Example of execution

```
mysqlfailover --discover-slaves-login=root:root \
--master=root:root@127.0.0.1:13001 auto
```

```
MySQL Replication Failover Utility
Failover Mode = auto      Next Interval = Fri Jan 31 09:49:17 2014

Master Information
-----
Binary Log File   Position  Binlog_Do_DB  Binlog_Ignore_DB
mysql-bin.000011  231

GTID Executed Set
453cdecc-82bd-11e3-9763-0800272864ba:1-4 [...]

Replication Health Status
+-----+-----+-----+-----+-----+-----+
| host      | port  | role   | state | gtid_mode | health |
+-----+-----+-----+-----+-----+-----+
| 127.0.0.1 | 13001 | MASTER | UP    | ON        | OK     |
| localhost | 13002 | SLAVE  | UP    | ON        | OK     |
| localhost | 13003 | SLAVE  | UP    | ON        | OK     |
+-----+-----+-----+-----+-----+-----+
```

If the master fails...

```
Failed to reconnect to the master after 3 attempts.  
  
Failover starting in 'auto' mode...  
# Candidate slave localhost:13002 will become the new master.  
# Checking slaves status (before failover).  
# Preparing candidate for failover.  
# Creating replication user if it does not exist.  
# Stopping slaves.  
# Performing STOP on all slaves.  
# Switching slaves to new master.  
# Disconnecting new master as slave.  
# Starting slaves.  
# Performing START on all slaves.  
# Checking slaves for errors.  
# Failover complete.  
# Discovering slaves for master at localhost:13002  
  
Failover console will restart in 5 seconds.
```

When failover is done

```
MySQL Replication Failover Utility
Failover Mode = auto      Next Interval = Fri Jan 31 09:54:52 2014

Master Information
-----
Binary Log File   Position  Binlog_Do_DB  Binlog_Ignore_DB
mysql-bin.000006  271

GTID Executed Set
04c3f4ae-89ba-11e3-84f4-0800272864ba:1 [...]

Replication Health Status
+-----+-----+-----+-----+-----+-----+
| host      | port  | role   | state | gtid_mode | health |
+-----+-----+-----+-----+-----+-----+
| localhost | 13002 | MASTER | UP    | ON        | OK     |
| localhost | 13003 | SLAVE  | UP    | ON        | OK     |
+-----+-----+-----+-----+-----+-----+
```

Limitations

- Monitoring node is not highly available
 - Closely monitor the monitoring node!
 - Manual failover with `mysqlrpladmin` may be preferred
- Errant transactions can prevent failover
 - Use `--pedantic` to get an error when starting `mysqlfailover`
 - Fix manually

Manual failover with mysqlrpladmin

- Planned promotion (switchover)

```
mysqlrpladmin --master=root:root@127.0.0.1:13002 \  
--new-master=root:root@127.0.0.1:13001 \  
--discover-slaves-login=root:root --demote-master \  
switchover
```

- Unplanned promotion (failover)

```
mysqlrpladmin  
--slaves=root:root@127.0.0.1:13002,root:root@127.0.0.  
1:13003 --candidates=root:root@localhost:13002  
failover
```

MySQL Fabric

- Available from MySQL Utilities 1.4
- Not limited to HA
- <http://www.mysql.fr/products/enterprise/fabric.html>
 - “Extensive framework for managing farms of MySQL servers”
- We had a previous webinar on Fabric
 - <http://www.percona.com/resources/mysql-webinars/putting-mysql-fabric-use>

Other solutions

- Older tools having added support for GTIDs
 - MHA from v0.56
 - Percona Replication Manager (PRM)
- However keep in mind that none of these tools checks errant transactions
 - Using them requires caution



PERCONA LIVE LONDON MYSQL CONFERENCE 2014

Learn from leading MySQL experts. **London, UK. November 3-4, 2014**

Early Bird Rates End August 31st
at 11:30pm BST

<http://www.percona.com/live/london-2014/>

Q&A

Thanks for attending!

stephane.combaudon@percona.com