



Troubleshooting Slow Queries

Sveta Smirnova
Principal Support Engineer
April, 28, 2016

Table of Contents

- Before we start
- What affects query execution
- EXPLAIN: how to find out how optimizer works
 - Data matters: again
- What really happened
 - Inside optimizer
 - Inside storage engine
 - Inside the server
- How to affect plans

Before we start

When you see slow query first

- You develop an application and find out that some queries are running slow
- After a while you find some slow queries in the slow query log
- All such queries are always slow
- We would not discuss cases when concurrency affects performance
 - This is subject of the next webinar

Slow is relative

- Mind you data!
- 75,000,000 rows
 - (INT, INT)
 - $75,000,000 * (4 + 4) = 600,000,000$ bytes = 572 MB
 - (INT, INT, DATETIME, VARCHAR(255), VARCHAR(255))
 - $75,000,000 * (4 + 4 + 8 + 256 + 256) = 39,600,000,000$ bytes = 37 G
 - $39,600,000,000 / 600,000,000 = 66$

Slow is relative


- Mind you data!
- Mind use case
 - Popular website
 - Admin interface
 - Weekly cron job

Slow is relative

- Mind you data!
- Mind use case
- Mind location
 - Server, used by multiple connections
 - Dedicated for OLAP queries


What affects query execution

Query execution workflow

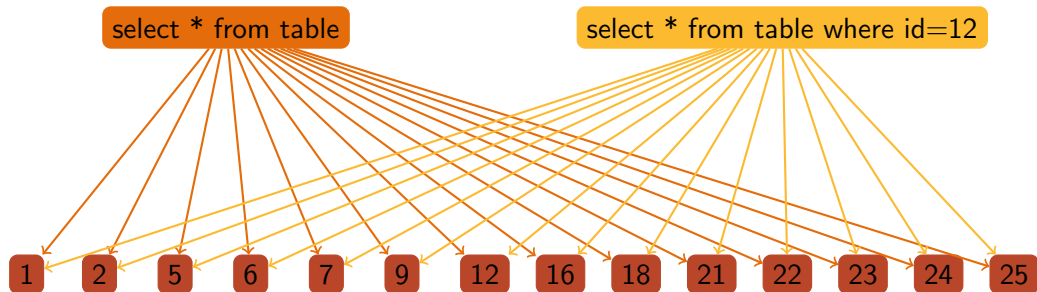


MySQL Indexes

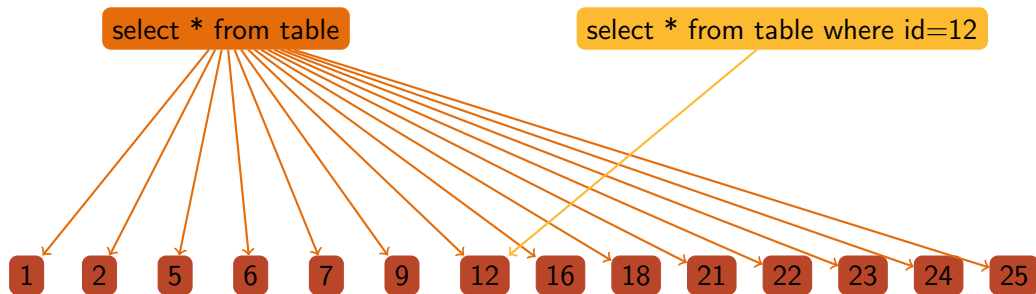
- B-Tree (Mostly)
- Fractal Tree
- R-Tree (Spatial)
- Hash (Memory SE)
- Engine-dependent



No index access



After index added



EXPLAIN: how to find out how optimizer works

MySQL EXPLAIN

- Estimates what happens during query execution
- Displays plan which Optimizer chooses
- Provides information how rows should be accessed

Effect of indexes: before

```
mysql> explain select * from t1\G
```

```
***** 1. row *****
```

```
...
```

```
rows: 12
```

```
Extra: NULL
```

```
mysql> explain select * from t1 where f2=12\G
```

```
***** 1. row *****
```

```
...
```

```
key: NULL
```

```
...
```

```
rows: 12
```

```
Extra: Using where
```

Same number of examined rows for both queries

Effect of indexes: after

```
mysql> alter table t1 add index(f2);  
Query OK, 12 rows affected (0.07 sec)  
Records: 12  Duplicates: 0  Warnings: 0
```

```
mysql> explain select * from t1 where f2=12\G  
***** 1. row *****  
  
...  
key: f2  
key_len: 5  
ref: const  
rows: 1  
Extra: NULL  
1 row in set (0.00 sec)
```

Much more effective!
Only 1 row examined

EXPLAIN: overview

mysql> explain extended select * from t1 join t2 where t1.int_key=1;

id	select_type	table	type	possible_keys	key	key_len	ref	rows	f...	Extra
1	SIMPLE	t1	ref	int_key,ik	int_key	5	const	4	100.	NULL
1	SIMPLE	t2	index	NULL	pk	9	NULL	6	100.	Using index

2 rows in set (0.00 sec) 1 warning (0.00 sec)

Note (Code 1003): /* select#1 */ select 'test'.'t1'.'pk' AS 'pk','test'.'t1'.'int_key' AS 'int_key','test'.'t2'.'pk' AS 'pk','test'.'t2'.'int_key' AS 'int_key' from 'test'.'t1' join 'test'.'t2' where ('test'.'t1'.'int_key' = 1)

Number of select

Tables, for which information is printed

Possible keys

Length of the key

Number of examined rows

Additional information

Table, for which information is printed

Product of rows here: how many rows in all tables will be accessed
For this example estimated value is $4 \times 6 = 24$

Select type

Key, which was actually used

% of filtered rows
rows x filtered / 100 — number of rows,
which will be joined with another table

How data is accessed

Which columns were compared with the index

Actual (optimized) query as executed by MySQL Server

EXPLAIN in details

```
mysql> explain extended select * from t1 join t2 where...
```

+---+-----+-----+-----+---				
id select_type table type ***				
+---+-----+-----+-----+---				
1	SIMPLE	t1	ref	***
1	SIMPLE	t2	index	***
+---+-----+-----+-----+---				

```
2 rows in set, 1 warning (0.00 sec)
```

SIMPLE;PRIMARY;UNION;DEPENDENT UNION;UNION RESULT;
SUBQUERY;DEPENDENT SUBQUERY;DERIVED;MATERIALIZED

system
const
eq_ref
ref
fulltext
ref_or_null
index_merge
unique_subquery
index_subquery
range
index
ALL

EXPLAIN in details: keys

Keys, which can be used for resolving the query

Actual length of the key (Important for multiple-column keys)

```
mysql> explain extended select * from t1 join t2 where t1.int_key=1;
```

```
***+-----+-----+-----+-----+***
***| possible_keys | key      | key_len | ref      |***
***+-----+-----+-----+-----+***
***| int_key, ik    | int_key  | 5        | const    |***
***| NULL           | pk       | 9        | NULL     |***
***+-----+-----+-----+-----+***
```

```
2 rows in set, 1 warning (0.00 sec)
```

Constant
Numeric in our case

Index used
to resolve rows

Only one key was actually used

Which columns were compared with the index

EXPLAIN in details: rows

Number of rows accessed

% of rows filtered

Additional information:
how query is resolved

- Using filesort
- Using temporary
- etc.

```
mysql> explain extended select * from t1 join t2 where t1.int_key=1;
```

```
***+-----+-----+-----+
***| rows  | filtered | Extra |
***+-----+-----+-----+
***| 4X6    |         |       |
***| =      | 100.00 | NULL  |
***| 24     | 100.00 | Using index; Using join buffer (Block Nested Loop) |
***+-----+-----+-----+
```

```
2 rows in set, 1 warning (0.00 sec)
```

All rows used

EXPLAIN type by example: ALL

```
mysql> explain select count(*) from employees where hire_date > '1995-01-01'\n\n***** 1. row *****\n      id: 1\n  select_type: SIMPLE\n        table: employees\n         type: ALL\npossible_keys: NULL\n          key: NULL\n      key_len: NULL\n         ref: NULL\n        rows: 300157\n  Extra: Using where\n1 row in set (0.00 sec)
```

All rows in the table examined
Worst plan ever!

EXPLAIN type by example: index

```
mysql> explain select count(*) from titles where title='Senior Engineer'\G
***** 1. row *****
```

```
      id: 1
  select_type: SIMPLE
        table: titles
      type: index
possible_keys: NULL
         key: emp_no
      key_len: 4
         ref: NULL
       rows: 444033
  Extra: Using where; Using index
```

No row in the table was accessed to resolve the query!
Only index used
Still all records in the index were scanned

```
1 row in set (0.11 sec)
```

EXPLAIN type by example: range

- We need to add index to table employees first
- ```
mysql> alter table employees add index(hire_date);
```

  
Query OK, 0 rows affected (3.48 sec)  
Records: 0 Duplicates: 0 Warnings: 0

# EXPLAIN type by example: range

```
mysql> explain select count(*) from employees where hire_date>'1995-01-01'\G
***** 1. row *****
```

id: 1

select\_type: SIMPLE

table: employees

type: range

possible\_keys: hire\_date

key: hire\_date

key\_len: 3

ref: NULL

rows: 68654

Extra: Using where; Using index

1 row in set (0.00 sec)

Only rows from given range used

Compare with ALL:  
 $300157 / 68654 = 4.3720$   
4 times less rows examined!



# EXPLAIN type by example: index\_subquery

```
mysql> explain select count(*), title from titles where from_date in
-> (select hire_date from employees where hire_date > '1995-01-01')
***** 1. row *****
 id: 1
 select_type: PRIMARY
 table: titles
 type: index
possible_keys: NULL
 key: emp_no
 key_len: 4
 ref: NULL
 rows: 443951
 Extra: Using where; Using index; Using temporary; Using filesort
```

# EXPLAIN type by example: index\_subquery

\*\*\*\*\* 2. row \*\*\*\*\*

id: 2

select\_type: DEPENDENT SUBQUERY

table: employees

type: index\_subquery

possible\_keys: hire\_date

key: hire\_date

key\_len: 3

ref: func

rows: 37

Extra: Using index; Using where

2 rows in set (0.00 sec)

# EXPLAIN type by example: unique\_subquery

```
mysql> explain select count(*), title from titles where emp_no in (select
-> emp_no from employees where hire_date > '1995-01-01') group by title\(
***** 1. row *****
 id: 1
 select_type: PRIMARY
 table: titles
 type: index
possible_keys: NULL
 key: emp_no
 key_len: 4
 ref: NULL
 rows: 443951
 Extra: Using where; Using index; Using temporary; Using filesort
```

# EXPLAIN type by example: unique\_subquery

\*\*\*\*\* 2. row \*\*\*\*\*

id: 2

select\_type: DEPENDENT SUBQUERY

table: employees

type: unique\_subquery

possible\_keys: PRIMARY,hire\_date

key: PRIMARY

key\_len: 4

ref: func

rows: 1

Extra: Using where

2 rows in set (0.00 sec)

# EXPLAIN type by example: index\_merge

- We need to modify table to run this test

```
mysql> create table dept_emp_copy (emp_no int, dept_no int);
Query OK, 0 rows affected (0.13 sec)
```

```
mysql> alter table dept_emp_copy add key(dept_no);
Query OK, 0 rows affected (1.32 sec)
Records: 0 Duplicates: 0 Warnings: 0
```

```
mysql> insert into dept_emp_copy(emp_no)
-> select distinct emp_no from dept_emp;
Query OK, 300024 rows affected (4.63 sec)
Records: 300024 Duplicates: 0 Warnings: 0
```

# EXPLAIN type by example: index\_merge

```
mysql> alter table dept_emp_copy add primary key(emp_no);
Query OK, 300024 rows affected (5.48 sec)
Records: 300024 Duplicates: 0 Warnings: 0
```

```
mysql> update dept_emp_copy, dept_emp
-> set dept_emp_copy.dept_no=dept_emp.dept_no
-> where dept_emp_copy.emp_no=dept_emp.emp_no;
Query OK, 300024 rows affected, 65535 warnings (15.66 sec)
Rows matched: 300024 Changed: 300024 Warnings: 0
```

# EXPLAIN type by example: index\_merge

```
mysql> explain select * from dept_emp_copy where
 -> dept_no > 5 or (emp_no > 10000 and emp_no < 20000)\G
***** 1. row *****
 id: 1
 select_type: SIMPLE
 table: dept_emp_copy
 type: index_merge
possible_keys: PRIMARY,dept_no
 key: dept_no,PRIMARY
 key_len: 5,4
 ref: NULL
 rows: 21103
 Extra: Using sort_union(dept_no,PRIMARY); Using where
```

# EXPLAIN type by example: original table

```
mysql> explain select * from dept_emp where dept_no in ('d005', 'd006',
-> 'd007', 'd008', 'd009') or (emp_no > 10000 and emp_no < 20000)\G
***** 1. row *****
 id: 1
select_type: SIMPLE
 table: dept_emp
 type: ALL
possible_keys: PRIMARY,emp_no,dept_no
 key: NULL
key_len: NULL
 ref: NULL
 rows: 332289
Extra: Using where
```



# EXPLAIN type by example: ref

```
mysql> explain select * from dept_emp where dept_no = 'd005'\G
***** 1. row *****
 id: 1
 select_type: SIMPLE
 table: dept_emp
 type: ref
possible_keys: dept_no
 key: dept_no
 key_len: 4
 ref: const
 rows: 145708
 Extra: Using where
1 row in set (0.00 sec)
```

# EXPLAIN type by example: eq\_ref

```
mysql> explain select * from dept_manager
-> join employees using(emp_no) limit 10\G
***** 1. row *****
 id: 1
select_type: SIMPLE
 table: dept_manager
 type: ALL
possible_keys: PRIMARY,emp_no
 key: NULL
 key_len: NULL
 ref: NULL
 rows: 24
 Extra:
```

# EXPLAIN type by example: eq\_ref

\*\*\*\*\* 2. row \*\*\*\*\*

id: 1

select\_type: SIMPLE

table: employees

type: eq\_ref

possible\_keys: PRIMARY

key: PRIMARY

key\_len: 4

ref: employees.dept\_manager.emp\_no

rows: 1

Extra:

2 rows in set (0.00 sec)

# EXPLAIN type by example: const

```
mysql> explain select * from departments where dept_no='d005'\G
***** 1. row *****
 id: 1
 select_type: SIMPLE
 table: departments
 type: const
possible_keys: PRIMARY
 key: PRIMARY
 key_len: 4
 ref: const
 rows: 1
 Extra:
1 row in set (0.00 sec)
```

# EXPLAIN PARTITIONS

```
mysql> explain partitions select count(*)
-> from employees_part where hire_date > '1991-01-01'\G
***** 1. row *****
 id: 1
select_type: SIMPLE
 table: employees_part
 partitions: p1,p2
 type: index
possible_keys: NULL
 key: PRIMARY
 key_len: 7
 ref: NULL
 rows: 135214
```

# EXPLAIN EXTENDED

```
mysql> explain extended select count(*) from employees join titles
 -> using(emp_no) where title='Senior Engineer'\G
```

```
***** 1. row *****
...
2 rows in set, 1 warning (0.00 sec)
```

```
mysql> show warnings\G
```

```
***** 1. row *****
Level: Note
Code: 1003
```

```
Message: select count(0) AS 'count(*)' from 'employees'.'employees' join
'employees'.'titles where (('employees'.'titles'.'emp_no' = 'employees'
.'employees'.'emp_no') and ('employees'.'titles'.'title' = 'Senior Engineer')
1 row in set (0.01 sec)
```

# EXPLAIN FORMAT = JSON

- More information than in regular EXPLAIN
  - Cost statistics

```
"duplicates_removal": {
 ...
 "cost_info": {
 "read_cost": "1252.00",
 "eval_cost": "88544.80",
 "prefix_cost": "89796.80",
 "data_read_per_join": "27M"
 },
```

# EXPLAIN FORMAT = JSON

- More information than in regular EXPLAIN
  - Cost statistics
  - Which part of index chosen

```
mysql> explain format=json SELECT first_name, last_name FROM employees
-> WHERE first_name='Steve' and last_name like 'V%'
-> and hire_date > '1990-01-01'\G
```

```
***** 1. row *****
EXPLAIN: {
```

```
...
```

```
"used_key_parts": [
 "first_name",
 "last_name"
],
```



# EXPLAIN FORMAT = JSON

- More information than in regular EXPLAIN
  - Cost statistics
  - Which part of index chosen
  - Columns, used to resolve query

```
mysql> explain format=json select count(*) from Country
 -> where Continent='Africa' and Population > 1000000\G
***** 1. row *****
...
 "used_columns": [
 "Continent",
 "Population"
],
```

# EXPLAIN FORMAT = JSON

- More information than in regular EXPLAIN
- Better structured view
  - Clear distinction for which of operations particular optimization used

```
mysql> explain format=json select distinct last_name
-> from employees order by last_name asc\G
```

...

```
"ordering_operation": {
 "using_filesort": false, - No temporary table here!
 "duplicates_removal": {
 "using_temporary_table": true,
 "using_filesort": true,
```

# EXPLAIN FORMAT = JSON

- More information than in regular EXPLAIN
- Better structured view
  - Clear distinction for which of operations particular optimization used
  - Easier to find out "which table belongs to which subselect" for complicated queries

# EXPLAIN FORMAT = JSON

- More information than in regular EXPLAIN
- Better structured view
  - Clear distinction for which of operations particular optimization used
  - Easier to find out "which table belongs to which subselect" for complicated queries
  - Separate member for each kind of optimization: grouping, ordering, duplicates\_removal, etc.

Data matters: again

# EXPLAIN: All, Index

- "All" always used for queries like `SELECT * FROM table;` without `WHERE` condition
- "Index" always used for queries like `SELECT indexed_field FROM table;` without `WHERE` condition
- Think carefully if you need to improve such queries, especially in cases when this is small table, joined with bigger one

# EXPLAIN: index\_merge, ref\_or\_null, ref

---

- index\_merge does not work with strings
- ref\_or\_null cannot become ref if you need NULL values
- ref cannot become eq\_ref if you need not unique values

# Queries limitations

- Queries which use LIKE '%foo%' cannot use indexes
- Function usage can prevent Optimizer to choose index access
  - Internal functions: with exceptions
  - Stored functions: always
  - UDF functions: always



# What really happened

# Inside optimizer

# Optimizer trace

- INFORMATION\_SCHEMA.OPTIMIZER\_TRACE
- Shows what happens inside Optimizer during query processing
- Typical use case

```
SET optimizer_trace="enabled=on";
SELECT ...;
SELECT * FROM INFORMATION_SCHEMA.OPTIMIZER_TRACE;
SET optimizer_trace="enabled=off";
```

# Optimizer trace example

- The table

```
mysql> show create table titles\G
```

```
***** 1. row *****
```

```
Table: titles
```

```
Create Table: CREATE TABLE 'titles' (
```

```
 'emp_no' int(11) NOT NULL,
```

```
 'title' varchar(50) NOT NULL,
```

```
 'from_date' date NOT NULL,
```

```
 'to_date' date DEFAULT NULL,
```

```
 PRIMARY KEY ('emp_no','title','from_date'),
```

```
 KEY 'emp_no' ('emp_no'),
```

```
 CONSTRAINT 'titles_ibfk_1' FOREIGN KEY ('emp_no') REFERENCES...
```

```
) ENGINE=InnoDB DEFAULT CHARSET=latin1
```

# Optimizer trace example

- The query

```
mysql> select distinct title from titles where year(from_date) > '1990';
```

```
+-----+
```

```
| title |
```

```
+-----+
```

```
| Staff |
```

```
| Senior Engineer |
```

```
...
```

# Optimizer trace example

- Possible and chosen keys

```
mysql> explain select distinct title from titles
```

```
-> where year(from_date) > '1990'\G
```

```
***** 1. row *****
```

```
id: 1
```

```
select_type: SIMPLE
```

```
table: titles
```

```
partitions: NULL
```

```
type: index
```

```
possible_keys: PRIMARY,emp_no
```

```
key: emp_no
```

```
key_len: 4
```

```
...
```

# Optimizer trace example

- EXPLAIN FORMAT=JSON sheds some light
  - Query cost for key emp\_no:

```
mysql> explain format=json select distinct title from titles
 -> where year(from_date) > '1990'\G
```

```
***** 1. row *****
```

```
EXPLAIN: {
 "query_block": {
 "select_id": 1,
 "cost_info": {
 "query_cost": "89630.20"
 },
 },
}
```

# Optimizer trace example

- EXPLAIN FORMAT=JSON sheds some light
  - Query cost for PRIMARY KEY:

```
mysql> explain format=json select distinct title from titles
 -> force index(primary) where year(from_date) > '1990'\G
***** 1. row *****
EXPLAIN: {
 "query_block": {
 "select_id": 1,
 "cost_info": {
 "query_cost": "530270.20"
 },
```



# Optimizer trace example

---

- Difference is  $(530270.20/89630.20)$  - 6 times, but why?

# Optimizer trace example

- Lets check optimizer trace
  - join\_optimization. rows\_estimation.  
potential\_range\_indexes

```
{
 "index": "PRIMARY",
 "usable": true,
 ...
},
{
 "index": "emp_no",
 "usable": true,
 ...
}
```

# Optimizer trace example

- Lets check optimizer trace
  - best\_covering\_index\_scan

```
"best_covering_index_scan": {
 "index": "emp_no",
 "cost": 91752,
 "chosen": false,
 "cause": "cost"
},
```

# Optimizer trace example

- Lets check optimizer trace
  - Chosen?

```
"group_index_range": {
 "distinct_query": true,
 "potential_group_range_indexes": [
 {
 "index": "PRIMARY",
 "covering": true,
 "usable": false,
 "cause": "select_attribute_not_prefix_in_index"
 },
```

# Optimizer trace example

- Lets check optimizer trace
  - Chosen?

```
{
 "index": "emp_no",
 "covering": true,
 "usable": false,
 "cause": "select_attribute_not_prefix_in_index"
}
```

# Optimizer trace example

- Lets check optimizer trace
  - Both indexes are not usable!

```
"best_access_path": {
 "considered_access_paths": [
 {
 "rows_to_scan": 441891,
 "access_type": "scan",
 "resulting_rows": 441891,
 "cost": 89630,
 "chosen": true
 }
]
},
```

# Optimizer trace example

- Lets check optimizer trace
  - emp\_no not used to resolve rows:

```
"best_covering_index_scan": {
 "index": "emp_no",
 "cost": 91752,
 "chosen": false,
 "cause": "cost"
},
```

# Optimizer trace example

- Lets check optimizer trace
  - So why we see emp\_no in EXPLAIN output?

```
"reconsidering_access_paths_for_index_ordering": {
 "clause": "GROUP BY",
 "index_order_summary": {
 "table": "'titles'",
 "index_provides_order": false,
 "order_direction": "undefined",
 "index": "emp_no",
 "plan_changed": false
 }
}
```



# ANALYZE in MariaDB

- Sometimes EXPLAIN lies

```
MariaDB [test]> explain select * from ol where thread_id=10432
-> and site_id != 9939 order by id limit 3\G
```

```
***** 1. row *****
```

|                          |  |                    |
|--------------------------|--|--------------------|
| id: 1                    |  | ref: NULL          |
| select_type: SIMPLE      |  | rows: 31           |
| table: ol                |  | Extra: Using where |
| type: index              |  |                    |
| possible_keys: thread_id |  |                    |
| key: PRIMARY             |  |                    |
| key_len: 4               |  |                    |
| 1 row in set (0.00 sec)  |  |                    |

# ANALYZE in MariaDB

- Sometimes EXPLAIN lies
- ANALYZE will show real numbers

```
MariaDB [test]> analyze select * from ol where thread_id=10432
-> and site_id != 9939 order by id limit 3\G
```

```
***** 1. row *****
```

|                          |  |                    |
|--------------------------|--|--------------------|
| id: 1                    |  | ref: const         |
| select_type: SIMPLE      |  | rows: 93283        |
| table: ol                |  | r_rows: 100000.00  |
| type: index              |  | filtered: 9.61     |
| possible_keys: thread_id |  | r_filtered: 0.00   |
| key: PRIMARY             |  | Extra: Using where |
| key_len: 4               |  |                    |

```
1 row in set (0.03 sec)
```

# Inside storage engine

# Handler\_\* status variables

- Same case as in previous example

```
mysql> explain select * from ol
```

```
-> where thread_id=10432 and site_id != 9939 order by id limit 3\G
```

```
***** 1. row *****
```

|                          |  |                    |
|--------------------------|--|--------------------|
| id: 1                    |  | ref: NULL          |
| select_type: SIMPLE      |  | rows: 33           |
| table: ol                |  | filtered: 8.07     |
| partitions: NULL         |  | Extra: Using where |
| type: index              |  |                    |
| possible_keys: thread_id |  |                    |
| key: PRIMARY             |  |                    |
| key_len: 4               |  |                    |

```
1 row in set, 1 warning (0,00 sec)
```

# Handler\_\* status variables

- Same case as in previous example
- Status variables 'Handler\_\*' will show truth

```
mysql> flush status; select * from ol
 -> where thread_id=10432 and site_id != 9939 order by id limit 3;
```

```
mysql> show status like 'Handler%';
```

| Variable_name      | Value  |
|--------------------|--------|
| ...                |        |
| Handler_read_first | 1      |
| Handler_read_key   | 1      |
| Handler_read_last  | 0      |
| Handler_read_next  | 100000 |

# Inside the server

# PROCESSLIST

---

- SHOW [FULL] PROCESSLIST
- INFORMATION\_SCHEMA.PROCESSLIST
- Your first alert in case of performance issue
- Shows all queries, run at the current moment

# Execution stages

- Can be seen in PROCESSLIST

```
mysql> show processlist\G
```

```
***** 1. row *****
```

```
Id: 7
```

```
User: root
```

```
Host: localhost:48799
```

```
db: employees
```

```
Command: Query
```

```
Time: 2
```

```
State: Sending data
```

```
Info: select count(*) from employees join titles using(emp_no)
 where title='Senior Engineer'
```

```
...
```



# Execution stages

---

- Can be seen in PROCESSLIST
  - Very useful when you need to answer on question:  
"What is my server doing now?"

# Execution stages

## • PERFORMANCE\_SCHEMA.EVENTS\_STAGES\_\*

```
mysql> select eshl.event_name, substr(sql_text, 1, 15) as 'sql',
-> eshl.timer_wait/1000000000000 w_s from events_stages_history_long
-> eshl join events_statements_history_long esthl on
-> (eshl.nesting_event_id = esthl.event_id) where
-> esthl.current_schema='employees' and sql_text like
-> 'select count(*) from employees%' order by eshl.timer_start asc;
```

| event_name                     | sql             | w_s    |
|--------------------------------|-----------------|--------|
| stage/sql/starting             | select count(*) | 0.0002 |
| stage/sql/checking permissions | select count(*) | 0.0000 |
| ...                            |                 |        |

# Execution stages

- PERFORMANCE\_SCHEMA.EVENTS\_STAGES\_\*

...

|                                |                 |        |  |
|--------------------------------|-----------------|--------|--|
| stage/sql/checking permissions | select count(*) | 0.0000 |  |
| stage/sql/Opening tables       | select count(*) | 0.0000 |  |
| stage/sql/init                 | select count(*) | 0.0001 |  |
| stage/sql/System lock          | select count(*) | 0.0000 |  |
| stage/sql/optimizing           | select count(*) | 0.0000 |  |
| stage/sql/statistics           | select count(*) | 0.0001 |  |
| stage/sql/preparing            | select count(*) | 0.0000 |  |
| stage/sql/executing            | select count(*) | 0.0000 |  |
| stage/sql/Sending data         | select count(*) | 5.4915 |  |
| stage/sql/end                  | select count(*) | 0.0000 |  |

...

# Temporary tables and other job

- Status variables

```
mysql> flush status;
```

```
Query OK, 0 rows affected (0,01 sec)
```

```
mysql> select count(*) from employees join titles using(emp_no)
 -> where title='Senior Engineer';
```

```
+-----+
```

```
| count(*) |
```

```
+-----+
```

```
| 97750 |
```

```
+-----+
```

```
1 row in set (5,44 sec)
```

# Temporary tables and other job

- Status variables

```
mysql> select * from performance_schema.session_status
-> where variable_name in ('Created_tmp_tables',
-> 'Created_tmp_disk_tables', 'Select_full_join',
-> 'Select_full_range_join', 'Select_range',
-> 'Select_range_check', 'Select_scan', 'Sort_merge_passes',
-> 'Sort_range', 'Sort_rows', 'Sort_scan') and variable_value > 0;
```

|   |               |   |                |   |
|---|---------------|---|----------------|---|
| + | -----         | + | -----          | + |
|   | VARIABLE_NAME |   | VARIABLE_VALUE |   |
| + | -----         | + | -----          | + |
|   | Select_scan   |   | 2              |   |
| + | -----         | + | -----          | + |

1 row in set (0,00 sec)

# Temporary tables and other job

- Status variables
- PERFORMANCE\_SCHEMA.EVENTS\_STATEMENTS\_\*

```
mysql> select * from performance_schema.events_statements_history_long
-> where sql_text like 'select count(*) from employees join %'\G
```

```
***** 1. row *****
```

```
...
```

|                            |                       |
|----------------------------|-----------------------|
| ROWS_SENT: 1               | SELECT_RANGE_CHECK: 0 |
| ROWS_EXAMINED: 541058      | SELECT_SCAN: 1        |
| CREATED_TMP_DISK_TABLES: 0 | SORT_MERGE_PASSES: 0  |
| CREATED_TMP_TABLES: 0      | SORT_RANGE: 0         |
| SELECT_FULL_JOIN: 0        | SORT_ROWS: 0          |
| SELECT_FULL_RANGE_JOIN: 0  | SORT_SCAN: 0          |
| SELECT_RANGE: 0            | NO_INDEX_USED: 0      |

# Temporary tables and other job

- Status variables
- PERFORMANCE\_SCHEMA.EVENTS\_STATEMENTS\_\*
- sys.statement\_analysis

```
mysql> select * from statement_analysis where query like 'SELECT COUNT
-> (*) FROM 'emplo%' and db='employees'\G
```

```
***** 1. row *****
 query: SELECT COUNT (*) FROM 'emplo ... 'emp_no') WHE...
 db: employees max_latency: 5.59 s
 full_scan: avg_latency: 5.41 s
exec_count: 7 lock_latency: 2.24 ms
 err_count: 0 rows_sent: 7
 warn_count: 0 rows_sent_avg: 1
total_latency: 37.89 s rows_examined: 3787406
```

# Temporary tables and other job

- Status variables
- PERFORMANCE\_SCHEMA.EVENTS\_STATEMENTS\_\*
- sys.statement\_analysis

rows\_examined\_avg: 541058

rows\_affected: 0

rows\_affected\_avg: 0

tmp\_tables: 0

tmp\_disk\_tables: 0

rows\_sorted: 0

sort\_merge\_passes: 0

digest: 4086bc3dc6510a1d9c8f2fe1f59f0943

first\_seen: 2016-04-14 15:19:19

last\_seen: 2016-04-14 16:13:14



# How to affect plans

# What affects optimizer plans?

---

- Index statistics
- Optimizer switches
- Bugs in optimizer

# Index statistics

---

- Collected by storage engine
- Optimizer decides which index to choose based on that

# Index statistics

- Can be viewed by SHOW INDEX command

```
mysql> show index from sbtest1;
```

| Table   | Key_name | Column_name | Cardinality |
|---------|----------|-------------|-------------|
| sbtest1 | k_1      | k           | 49142       |

```
mysql> select count(distinct id), count(distinct k) from sbtest1;
```

| count(distinct id) | count(distinct k) |
|--------------------|-------------------|
| 100000             | 17598             |

# Index statistics

- Can be updated
  - ANALYZE TABLE
  - If does not help: rebuild table
    - OPTIMIZE TABLE
    - ALTER TABLE ENGINE=INNODB; ANALYZE TABLE

# Optimizer switches

```
mysql> select @@optimizer_switch\G
***** 1. row *****
@@optimizer_switch: index_merge=on,index_merge_union=on,
index_merge_sort_union=on,index_merge_intersection=on,
engine_condition_pushdown=on,index_condition_pushdown=on,
mrr=on,mrr_cost_based=on,
block_nested_loop=on,batched_key_access=off,
materialization=on,semijoin=on,loosescan=on,firstmatch=on,
duplicateweedout=on,subquery_materialization_cost_based=on,
use_index_extensions=on,condition_fanout_filter=on,derived_merge=on
1 row in set (0,00 sec)
```

# Optimizer switches

- Turn ON and OFF particular optimization
- Can be not helpful
  - Especially for queries, tuned for previous versions
- Work with them as with any other option
  - Turn OFF and try

```
SET optimizer_switch = 'use_index_extensions=off';
```

```
SELECT ...
```

```
EXPLAIN SELECT ...
```

- If helps implement in queries

```
SELECT /*+ SEMIJOIN(FIRSTMATCH, LOOJESCAN) */ * FROM t1 ...;
```

```
SELECT /*+ BKA(t1) NO_BKA(t2) */ * FROM t1 INNER JOIN t2 WHERE ...;
```

# Bugs in Optimizer

- Optimizer choses wrong index for no reason
- Statistics is up to date
- Solution
  - Use index hints
    - FORCE INDEX
    - IGNORE INDEX
- On every upgrade
  - Remove index hints
  - Test if query improved
  - You must do it even for minor version upgrades!



# Summary

# Summary

---

- Understanding EXPLAIN output is essential for query tuning
- But real job is done by storage engine
- Index statistics affect query execution plan
- All index hints, optimizer hints and other workarounds must be validated on each upgrade

# More information

---

- EXPLAIN Syntax
- EXPLAIN FORMAT=JSON is Cool! series
- Troubleshooting Performance add-ons
- Optimizer Hints
- Custom Hint Plugin

# Place for your questions

---

???

# Thank you!

---

`http://www.slideshare.net/SvetaSmirnova`

`https://twitter.com/svetsmirnova`