



Introduction to troubleshooting How to create test setup

Sveta Smirnova
Principal Support Engineer
March, 2, 2016

Table of Contents

- Before troubleshooting begins
- Test setup
 - First considerations
 - Simplest setup
 - Two kinds of issues with databases
 - Avoiding limitations

**Before
troubleshooting
begins**

Importance of tests

- When something goes wrong it is often not clear what causes the issue
- Best way to clarify it is to repeat the problem
 - You can set breakpoints
 - You can experiment with input
 - You can rewrite query or code

Can we ignore tests?

- Experienced users sometimes can know what leads to one or another issue
 - Because they repeated it in their environments dozens of times!
- Not all issues can be solved without testing
 - Brand new bugs with new features
 - Issues, badly affected by data
 - Most of issues!
- I myself test something new everyday

Troubleshooting webinars

- In future webinars I would often ask you to repeat something
- In all such cases I would expect you have properly set test environment
- Tests on production are dangerous
 - Big data uses a lot of resources, required by your application
 - Repeating crashing issues can kill production server

Troubleshooting webinars

- In future webinars I would often ask you to repeat something
- In all such cases I would expect you have properly set test environment
- Tests on production are dangerous
- How to create test environment more effectively?
 - This is the topic of our today webinar.

Test setup

First considerations

Why don't use development server?

- When working on an application we usually first develop it on our own machine
- In case of database application we also usually have small database with test data

Table : User

Id	Name
1	test
• 2	Foo
...	
100000	Bar

- Usually has much more data

[illegible]

Production database

- Usually has much more data
- Data is not ordered and depends from random factors
- Can have extremely powerful hardware which costs too much to be duplicated on development server

Simplest setup

Ideal case

- If you can
 - Use same hardware as on production server for your test setup
 - Can copy data without modifications
 - Can make full binary backup of production server

Ideal case

- If you can
- Install same version as on production
 - Just install same package
 - Use MySQL Sandbox
 - On UNIX, Linux, Mac OS X

```
perl Makefile.PL
make
[make test] - optionally
sudo make install
make_sandbox your_MySQL/Percona_server.tar.gz
make_replication_sandbox your_MySQL/Percona_server.tar.gz
```

Ideal case

- If you can
- Install same version as on production
- Make backup and restore
 - Make full binary backup of your production server
 - `innobackupex /data/backups`
 - `cp -R /path/to/datadir /data/backups/datadir`
 - Restore it on the test server
 - Copy option file from production to test server
 - Start and use test server

We live in not ideal world

- You can have less power server for tests
 - Less space
 - Smaller number of CPU
 - Smaller amount of RAM
 - Have one test server and master-slave issue on production
 - Have other limitations
- There are ways to workaround these limitations

Still we need to have in mind the ideal

- MySQL/Percona/MariaDB Server must use same version as production
 - There is no excuse to test on different version!
- Same options
 - Not always possible
- Same data
 - Not always possible

Two kinds of issues with databases

Two groups of database issues

- Wrong behavior
 - You get results which you don't expect
 - Wrong rows updated
 - Database crashes
 - Slave has dataset different from master's
 - Etc.
- Performance issues

Two groups of database issues

- Wrong behavior
- Performance issues
 - Query runs slowly
 - Overall server performance is slower than you need
 - Queries hang up from time to time
 - Etc.

Two groups of database issues

- Wrong behavior
 - Most, but not all, issues can be repeated on development server
- Performance issues
 - Most issues can **not** be repeated on development server

Avoiding limitations

Lack of disk space

- Copy only relevant data
 - Only tables which participate in the problematic query
 - Part of data if table is huge
 - LIMIT
 - WHERE
 - mysqldump --where
 - Only columns which participate in the query
- Not always leads to issue repetition!

Sensitive data

- Copy only columns which do not have sensitive data
 - Example
 - Id, dates, public article content - all these are either generic or public already
 - Email, password, names - these can be sensitive
- Encrypt sensitive data
 - Non-symmetric encryption functions, such as md5, sha
 - Use substr and concat to have values of length, similar to data you have
- Replace sensitive data with random values

Lack of CPU cores

- Some options must be scaled
 - innodb_thread_concurrency
 - innodb_concurrency_tickets
 - If you changed innodb_thread_concurrency, check user manual
 - innodb_commit_concurrency
 - innodb_purge_threads
 - innodb_read_io_threads
 - innodb_write_io_threads
 - thread_concurrency
 - thread_cache_size
 - Only if you scale number of connections in your test

Lack of CPU cores

- Some options must be scaled
- Use scale coefficient: $\text{SCALE} = \frac{\text{NUMBER_OF_CPU_CORES_ON_PRODUCTION}}{\text{NUMBER_OF_CPU_CORES_ON_TEST}}$
 - $\text{innodb_thread_concurrency_test} = \frac{\text{innodb_thread_concurrency_production}}{\text{SCALE}}$
 - Makes sense only if such options are greater than number of CPU cores on the test server or unlimited

Scale coefficient example

Table : innodb_thread_concurrency scale

Production: 16 cores	Test: 4 cores	Adjust number of client threads?
0	0	It depends
4	4	Not
16	4	Yes
32	8	Yes
128	32	Yes

Disk-related options

- innodb_io_capacity
 - Depends on disk speed!
 - Can give very different results if disk speed is not taken in account
- innodb_io_capacity_max

Lack of RAM

- Some options must be scaled
- Use scale coefficient: $SCALE = \frac{RAM_ON_PRODUCTION}{RAM_ON_TEST}$
- Scale expected return time for performance tests as well
- You can use less data to repeat similar slowdown as on production
 - Care about data distribution!
 - Check EXPLAIN output
 - More waiting time can be better for testing: easier to notice difference

RAM-related options: scale them

- `innodb_buffer_pool_size`
- `innodb_log_file_size`
 - only if changed buffer pool size
- `innodb_buffer_pool_chunk_size`
 - only if very large on production, usually not needed
 - It is unlikely what you will have to tune InnoDB Buffer Pool resizing

RAM-related options: scale them

- innodb_buffer_pool_size
- innodb_log_file_size
- innodb_buffer_pool_chunk_size
- innodb_buffer_pool_instances
 - only if InnoDB Buffer Pool is very large and current number of instances does not make sense on test server.
 - Example: 96G buffer pool on production with 12 instances and you test on laptop with 8G memory

RAM-related options: scale them

- innodb_buffer_pool_size
- innodb_log_file_size
- innodb_buffer_pool_chunk_size
- innodb_buffer_pool_instances
- query_cache_size, key_buffer_size
 - only if larger than available memory
- Do not adjust session and operation-specific options unless they set to very large values

RAM-related options: scale them

- innodb_buffer_pool_size
- innodb_log_file_size
- innodb_buffer_pool_chunk_size
- innodb_buffer_pool_instances
- query_cache_size, key_buffer_size
- Do not adjust session and operation-specific options unless they set to very large values
- table_open_cache, table_definition_cache
 - Do not adjust

Session options: do not scale

- join_buffer_size
- [max_]binlog_[stmt_]cache_size
- max_heap_table_size
- tmp_table_size
- myisam_*
- net_buffer_length
- parser_max_mem_size
- read_buffer_size
- read_rnd_buffer_size
- sort_buffer_size

If you have to test on production

- Do not test crashes and issues which lead to full server hang!
- Create copy of database
 - CREATE DATABASE test_copy;
 - Copy all involved objects (tables, views and so on) into new database
 - Use any backup-restore technique
 - mysqldump test |mysql test_copy
 - mysqldbcopy --source=root:@127.0.0.1:13000 --destination=root:@127.0.0.1:13000 test:test_copy --multiprocess=4
 - Your favorite method

If you have to test on production

- Do not test crashes and issues which lead to full server hang!
- Create copy of database
- Create copy of single table

```
CREATE TABLE test_copy LIKE test;  
INSERT INTO test_copy SELECT * FROM test;  
INSERT INTO test_copy SELECT * FROM test WHERE condition;  
-- Useful when your query returns wrong results  
INSERT INTO test_copy SELECT * FROM test LIMIT 1000000;  
-- Speeds up performance tests on big data
```

Any backup-restore technique which can copy-restore single table

Summary

Summary

- There is no excuse to test on different version!
- 100 % copy of data and options is ideal
- Ideal does not always happen in real life
- Care about
 - Data
 - Options
- Use test environment for experiments

References

- MySQL Sandbox
- Percona XtraBackup
- mysqldump
- mysqldbcopy
- MySQL option tables

Place for your questions

???

Thank you!

`http://www.slideshare.net/SvetaSmirnova`

`https://twitter.com/svetsmirnova`